

# Get Started, Part 5: Stacks

*Estimated reading time: 10 minutes*

1: Orientation (<https://docs.docker.com/get-started/part1>)

2: Containers (<https://docs.docker.com/get-started/part2>)

3: Services (<https://docs.docker.com/get-started/part3>)

4: Swarms (<https://docs.docker.com/get-started/part4>)

5: Stacks (<https://docs.docker.com/get-started/part5>)

6: Deploy your app (<https://docs.docker.com/get-started/part6>)

## Prerequisites

- Install Docker version 1.13 or higher (<https://docs.docker.com/engine/installation/>).
- Get Docker Compose (<https://docs.docker.com/compose/overview/>) as described in Part 3 prerequisites (<https://docs.docker.com/get-started/part3/#prerequisites>).
- Get Docker Machine (<https://docs.docker.com/machine/overview/>) as described in Part 4 prerequisites (<https://docs.docker.com/get-started/part4/#prerequisites>).
- Read the orientation in Part 1 (<https://docs.docker.com/get-started/>).
- Learn how to create containers in Part 2 (<https://docs.docker.com/get-started/part2/>).
- Make sure you have published the `friendlyhello` image you created by pushing it to a registry (<https://docs.docker.com/get-started/part2/#share-your-image>). We use that shared image here.
- Be sure your image works as a deployed container. Run this command, slotting in your info for `username` , `repo` , and `tag` : `docker run -p 80:80 username/repo:tag` , then visit `http://localhost/` .
- Have a copy of your `docker-compose.yml` from Part 3 (<https://docs.docker.com/get-started/part3/>) handy.
- Make sure that the machines you set up in part 4 (<https://docs.docker.com/get-started/part4/>) are running and ready. Run `docker-machine ls` to verify this. If the machines are stopped, run `docker-machine start myvm1` to boot the manager, followed by `docker-machine start myvm2` to boot the worker.

- Have the swarm you created in part 4 (<https://docs.docker.com/get-started/part4/>) running and ready. Run `docker-machine ssh myvm1 "docker node ls"` to verify this. If the swarm is up, both nodes report a `ready` status. If not, reinitialize the swarm and join the worker as described in Set up your swarm (<https://docs.docker.com/get-started/part4/#set-up-your-swarm>).

## Introduction

In part 4 (<https://docs.docker.com/get-started/part4/>), you learned how to set up a swarm, which is a cluster of machines running Docker, and deployed an application to it, with containers running in concert on multiple machines.

Here in part 5, you reach the top of the hierarchy of distributed applications: the **stack**. A stack is a group of interrelated services that share dependencies, and can be orchestrated and scaled together. A single stack is capable of defining and coordinating the functionality of an entire application (though very complex applications may want to use multiple stacks).

Some good news is, you have technically been working with stacks since part 3, when you created a Compose file and used `docker stack deploy`. But that was a single service stack running on a single host, which is not usually what takes place in production. Here, you can take what you've learned, make multiple services relate to each other, and run them on multiple machines.

You're doing great, this is the home stretch!

## Add a new service and redeploy

It's easy to add services to our `docker-compose.yml` file. First, let's add a free visualizer service that lets us look at how our swarm is scheduling containers.

1. Open up `docker-compose.yml` in an editor and replace its contents with the following. Be sure to replace `username/repo:tag` with your image details.

```

version: "3"
services:
  web:
    # replace username/repo:tag with your name and image details
    image: username/repo:tag
    deploy:
      replicas: 5
      restart_policy:
        condition: on-failure
      resources:
        limits:
          cpus: "0.1"
          memory: 50M
    ports:
      - "80:80"
    networks:
      - webnet
  visualizer:
    image: dockersamples/visualizer:stable
    ports:
      - "8080:8080"
    volumes:
      - "/var/run/docker.sock:/var/run/docker.sock"
    deploy:
      placement:
        constraints: [node.role == manager]
    networks:
      - webnet
networks:
  webnet:

```

The only thing new here is the peer service to `web` , named `visualizer` . Notice two new things here: a `volumes` key, giving the visualizer access to the host's socket file for Docker, and a `placement` key, ensuring that this service only ever runs on a swarm manager -- never a worker. That's because this container, built from an open source project created by Docker (<https://github.com/ManoMarks/docker-swarm-visualizer>), displays Docker services running on a swarm in a diagram.

We talk more about placement constraints and volumes in a moment.

2. Make sure your shell is configured to talk to `myvm1` (full examples are here (<https://docs.docker.com/get-started/part4/#configure-a-docker-machine-shell-to-the-swarm-manager>)).

- Run `docker-machine ls` to list machines and make sure you are connected to `myvm1` , as indicated by an asterisk next to it.
- If needed, re-run `docker-machine env myvm1` , then run the given command to configure the shell.

On **Mac or Linux** the command is:

```
eval $(docker-machine env myvm1)
```

On **Windows** the command is:

```
& "C:\Program Files\Docker\Docker\Resources\bin\docker-machine.exe" en
```



3. Re-run the `docker stack deploy` command on the manager, and whatever services need updating are updated:

```
$ docker stack deploy -c docker-compose.yml getstartedlab
Updating service getstartedlab_web (id: angi1bf5e4to03qu9f93trnxm)
Creating service getstartedlab_visualizer (id: l9mnwkeq2jiononb5ihz9u7a4)
```



4. Take a look at the visualizer.

You saw in the Compose file that `visualizer` runs on port 8080. Get the IP address of one of your nodes by running `docker-machine ls`. Go to either IP address at port 8080 and you can see the visualizer running:

The screenshot shows the Docker Visualizer interface in a web browser. The browser's address bar displays the URL `192.168.99.101:8080`. The Visualizer interface has a dark blue background with a white Docker logo at the top. Below the logo, there are two main sections representing Docker stacks:

- myvm1 manager 1G RAM**: This stack contains three containers:
  - getstartedlab\_web**: image : get-started:part1@sha256:0788ab7acd921a2b519529efec, tag : part1@sha256:7799210bc, updated : 20/3 0:47, state : running
  - getstartedlab\_visualizer**: image : visualizer:stable@sha256:68d25fde535ca94ca75ac3e95a, tag : stable@sha256:f924ad66c, updated : 20/3 0:47, state : running
  - getstartedlab\_web**: image : get-started:part1@sha256:819db4416425a25889345aed8, tag : part1@sha256:7799210bc, updated : 20/3 0:47, state : running
- myvm2 worker 1G RAM**: This stack contains three containers:
  - getstartedlab\_web**: image : get-started:part1@sha256:029e9249ce0b94141f494b824, tag : part1@sha256:7799210bc, updated : 20/3 0:47, state : running
  - getstartedlab\_web**: image : get-started:part1@sha256:8d21a3c68f626b6c82f1e9ee62, tag : part1@sha256:7799210bc, updated : 20/3 0:47, state : running
  - getstartedlab\_web**: image : get-started:part1@sha256:6f8009d5f8dbb11ed1c994de7f, tag : part1@sha256:7799210bc, updated : 20/3 0:47, state : running

The single copy of `visualizer` is running on the manager as you expect, and the 5 instances of `web` are spread out across the swarm. You can corroborate this visualization by running `docker stack ps <stack>` :

```
docker stack ps getstartedlab
```

The visualizer is a standalone service that can run in any app that includes it in the stack. It doesn't depend on anything else. Now let's create a service that *does* have a dependency: the Redis service that provides a visitor counter.

## Persist the data

Let's go through the same workflow once more to add a Redis database for storing app data.

1. Save this new `docker-compose.yml` file, which finally adds a Redis service. Be sure to replace `username/repo:tag` with your image details.

```
version: "3"
services:
  web:
    # replace username/repo:tag with your name and image details
    image: username/repo:tag
    deploy:
      replicas: 5
      restart_policy:
        condition: on-failure
      resources:
        limits:
          cpus: "0.1"
          memory: 50M
    ports:
      - "80:80"
    networks:
      - webnet
  visualizer:
    image: dockersamples/visualizer:stable
    ports:
      - "8080:8080"
    volumes:
      - "/var/run/docker.sock:/var/run/docker.sock"
    deploy:
      placement:
        constraints: [node.role == manager]
    networks:
      - webnet
  redis:
    image: redis
    ports:
      - "6379:6379"
    volumes:
      - "/home/docker/data:/data"
    deploy:
      placement:
        constraints: [node.role == manager]
    command: redis-server --appendonly yes
    networks:
      - webnet
networks:
  webnet:
```

Redis has an official image in the Docker library and has been granted the short `image` name of just `redis`, so no `username/repo` notation here. The Redis port, 6379, has been pre-configured by Redis to be exposed from the container to the host, and here in our Compose file we expose it from the host to the world, so you can actually enter the IP for any of your nodes into Redis Desktop Manager and manage this Redis instance, if you so choose.

Most importantly, there are a couple of things in the `redis` specification that make data persist between deployments of this stack:

- `redis` always runs on the manager, so it's always using the same filesystem.
- `redis` accesses an arbitrary directory in the host's file system as `/data` inside the container, which is where Redis stores data.

Together, this is creating a "source of truth" in your host's physical filesystem for the Redis data. Without this, Redis would store its data in `/data` inside the container's filesystem, which would get wiped out if that container were ever redeployed.

This source of truth has two components:

- The placement constraint you put on the Redis service, ensuring that it always uses the same host.
- The volume you created that lets the container access `./data` (on the host) as `/data` (inside the Redis container). While containers come and go, the files stored on `./data` on the specified host persists, enabling continuity.

You are ready to deploy your new Redis-using stack.

2. Create a `./data` directory on the manager:

```
docker-machine ssh myvm1 "mkdir ./data"
```

3. Make sure your shell is configured to talk to `myvm1` (full examples are here (<https://docs.docker.com/get-started/part4/#configure-a-docker-machine-shell-to-the-swarm-manager>)).

- Run `docker-machine ls` to list machines and make sure you are connected to `myvm1`, as indicated by an asterisk next to it.
- If needed, re-run `docker-machine env myvm1`, then run the given command to configure the shell.

On **Mac or Linux** the command is:

```
eval $(docker-machine env myvm1)
```

On **Windows** the command is:

```
& "C:\Program Files\Docker\Docker\Resources\bin\docker-machine.exe" en
```



4. Run `docker stack deploy` one more time.




```
$ docker stack deploy -c docker-compose.yml getstartedlab
```

5. Run `docker service ls` to verify that the three services are running as expected.

```
$ docker service ls
```

ID	NAME	MODE	REPLICAS
x7uij6xb4foj	getstartedlab_redis	replicated	1/1
n5rvhm52ykq7	getstartedlab_visualizer	replicated	1/1
mifd433bti1d	getstartedlab_web	replicated	5/5

6. Check the web page at one of your nodes, such as <http://192.168.99.101>, and take a look at the results of the visitor counter, which is now live and storing information on Redis.



Also, check the visualizer at port 8080 on either node's IP address, and notice the `redis` service running along with the `web` and `visualizer` services.



The screenshot shows the Docker Visualizer interface in a web browser. The address bar displays `192.168.99.100:8080`. The interface features a dark blue background with a white Docker logo at the top. Below the logo, there are two main sections for container stacks. The left stack is titled **myvm1 manager 1G RAM** and the right stack is titled **myvm2 worker 1G RAM**. Each stack contains several container cards. The cards for **getstartedlab\_visualizer** and **getstartedlab\_redis** are highlighted with yellow and pink borders, respectively. The **getstartedlab\_web** cards have green borders. Each card displays the following information: container name, image, tag, update time, ID, and state (running).

Container Name	Image	Tag	Updated	ID	State
getstartedlab_visualizer	visualizer:stable	sha256:f92...	2/6 16:3	87ffd7636236cf7d4465e823711c692	running
getstartedlab_web	get-started:part1	sha256:0...	2/6 16:3	dd4cf1a1c3f2dd6aaf1411205f96b03	running
getstartedlab_redis	redis:latest	sha256:548a75066f3f28	2/6 16:3	809cfb25a8fbc3426b67d6f8da69336	running
getstartedlab_web	get-started:part1	sha256:0601c866aab2a	2/6 16:3	dbc3451300ae1f2489212f4e79e622	running
getstartedlab_web	get-started:part1	sha256:0601c866aab2a	2/6 16:3	e1ba272633d6915263fdb2a04934b1	running
getstartedlab_web	get-started:part1	sha256:0601c866aab2a	2/6 16:3	cc770a2a2a3bba2fb2a0961084f595	running

On to Part 6 >> (<https://docs.docker.com/get-started/part6/>)

## Recap (optional)

Here's a terminal recording of what was covered on this page (<https://asciinema.org/a/113840>):



```
ports:
  - "8080:8080"
volumes:
  - "/var/run/docker.sock:/var/run/docker.sock"
deploy:
  placement:
    constraints: [node.role == manager]
networks:
  - webnet
redis:
  image: redis
  ports:
    - "6379:6379"
  volumes:
    - ./data:/data
  deploy:
    placement:
      constraints: [node.role == manager]
  networks:
    - webnet
networks:
  webnet:
```

You learned that stacks are inter-related services all running in concert, and that -- surprise! -- you've been using stacks since part three of this tutorial. You learned that to add more services to your stack, you insert them in your Compose file. Finally, you learned that by using a combination of placement constraints and volumes you can create a permanent home for persisting data, so that your app's data survives when the container is torn down and redeployed.

stack (<https://docs.docker.com/glossary/?term=stack>), data (<https://docs.docker.com/glossary/?term=data>), persist (<https://docs.docker.com/glossary/?term=persist>), dependencies (<https://docs.docker.com/glossary/?term=dependencies>), redis (<https://docs.docker.com/glossary/?term=redis>), storage (<https://docs.docker.com/glossary/?term=storage>), volume (<https://docs.docker.com/glossary/?term=volume>), port (<https://docs.docker.com/glossary/?term=port>)